

A Unified Deep Learning Framework for Protein Stability, Interaction, and Expression

Md Ahbab Hamid Khan | 011223869

Ahnaf Atique — 0112230597

Shadhin Nandi — 0112230604

Md Wali Ullah Khan — 0112230754

Jamiul Hasan — 0112230871

Akib Bari — 0112231020

Department of Computer Science & Engineering

United International University

`mkhan223869@bscse.uiu.ac.bd`

April 4, 2026

1 Introduction

Accurate prediction of protein structure and function from amino acid sequences is one of the most enduring and complex challenges in computational biology. For decades, the central dogma of molecular biology – sequence determines structure and structure determines function – has driven the development of computational methods aimed at solving the “protein folding problem.” The recent emergence of AlphaFold2 (AF2) and its successor, AlphaFold3 (AF3), has been hailed as a definitive solution to the structural prediction component of this dogma [1–3]. These deep learning systems, using vast evolutionary histories encoded in Multiple Sequence Alignments (MSAs) and sophisticated attention mechanisms, have achieved an atomic-level accuracy that rivals experimental methods such as X-ray crystallography and cryo-electron microscopy (Cryo-EM) [1]. However, as the initial euphoria surrounding these breakthroughs settles, the field confronts a more nuanced and perhaps more difficult reality: biology is not static.

Proteins are not rigid sculptures frozen in space; they are dynamic molecular machines that traverse complex, rugged energy landscapes. Their function, whether it be enzymatic catalysis, signal transduction, or structural support, is inherently tied to their thermodynamic stability ($\Delta\Delta G$), their ability to adopt multiple conformational states, and their ability to interact with a diverse array of ligands, nucleic acids, and other proteins [3]. The current state-of-the-art (SOTA) models, while exceptional at predicting the native state of folded domains, often fail to capture these dynamic properties. They frequently hallucinate structures for intrinsically Disordered Regions (IDRs), struggle to predict the thermodynamic impact of single-point mutations without bias, and exhibit “fragile” generalisation when applied to sequences that diverge significantly from the evolutionary clusters seen during training [2, 4, 16].

1.1 The Limitations of the Static Paradigm

Reliance on static structural prediction creates a significant bottleneck for translational applications such as de novo protein design and variant pathogenicity interpretation. For example, determining whether a specific mutation in a patient’s genome is benign or pathogenic often hinges on subtle changes in protein stability or the perturbation of a disordered interaction motif—phenomena that purely structural models such as AF3 are not equipped to quantify [13–15]. As

noted in recent literature, AF3 creates a “static picture” that buttonholes the system, failing to model the essential molecular motions that accompany drug attachment or biological function [3]. Furthermore, the model’s reliance on MSAs renders it less effective for “orphan proteins”—sequences with few or no known homologs—thereby excluding a vast universe of potential synthetic proteins and rapidly evolving viral antigens from accurate analysis [8].

Moreover, the “black box” nature of these massive neural networks raises concerns about interpretability and accessibility. The computational burden required to run models like AF3 or train large-scale transformers restricts their utility to well-funded research centres, creating a disparity in scientific progress [1]. There is an urgent need for models that are not only accurate but also computationally efficient and mechanistically interpretable.

1.2 The Data and Generalisation Crisis

Parallel to the modelling challenge is the issue of data scarcity and heterogeneity. While structural databases like the PDB have grown, they pale in comparison to the sequence space. In domains such as mass spectrometry-based proteomics, the lack of standardised, large-scale datasets has hampered the development of robust machine learning tools [7, 12]. Models trained on limited datasets frequently overfit to local sequence patterns, failing to learn the underlying physicochemical rules—such as mRNA folding energy or codon usage bias—that govern protein expression and stability across different biological contexts [6]. The field requires a shift from “raw data” learning to “mechanistically informed” learning, where biological priors are integrated into the architecture to ensure robustness.

1.3 Objectives of This Work

This paper introduces a novel, unified machine learning framework designed to bridge the gap between static structure prediction and dynamic functional characterisation. Drawing upon a rigorous synthesis of recent advancements in Evolutionary Computation (EC), Geometric Deep Learning (GDL), and Protein Language Models (PLMs), we propose a Holo-GNN (Holographic Geometric Neural Network) architecture.

Our approach is distinct in three key areas:

- **Geometric Stacking of Mechanistic Features:** We move beyond simple one-hot encoding or raw PLM embeddings by explicitly integrating mechanistic features—such as mRNA thermodynamics and local energetic profiles—into the graph topology of the network. This ensures that the model learns biologically valid rules that generalise to unseen sequences [6].
- **Unified Stability and Dynamics Prediction:** We incorporate an antisymmetric Siamese network architecture to predict thermodynamic stability ($\Delta\Delta G$) without the bias common to older methods [9], coupled with a probabilistic output head capable of quantifying the ensemble dimensions of disordered regions [16].
- **Efficiency and Accessibility:** By leveraging pre-trained PLMs (like ESM-2) to extract evolutionary signal without the need for costly runtime MSA construction, our model offers high accuracy for orphan proteins and rapid inference speeds suitable for proteome-wide scanning [8, 18].

In the following sections, we provide an exhaustive review of the relevant literature, analysing the strengths and weaknesses of current methodologies. We then detail our methodology, specifying the provenance of our multi-species training datasets and articulating the novel contributions of our architecture to the field of computational proteomics.

2 Literature Review

The landscape of computational biology has been reshaped by the influx of deep learning techniques. This review categorises the literature into five thematic pillars: The Evolution of Structure Prediction; The Challenge of Thermodynamic Stability; The Unstructured Proteome; Sequence-to-Function Generalisation; and The Proteomics Data Ecosystem.

2.1 The Evolution of Structure Prediction: From Templates to Diffusion

The journey of Protein Structure Prediction (PSP) has traditionally been dichotomised into Template-Based Modelling (TBM) and Template-Free Modelling (TFM). TBM, including homology modelling and threading, relies on the detection of evolutionary relatives with solved structures. While historically accurate for conserved families, TBM fails for targets with low sequence identity to known structures [1]. TFM, or *ab initio* modelling, attempts to simulate the folding process from physical principles. Early EC-based methods treated this as a global optimisation problem, utilising Genetic Algorithms (GA) or Particle Swarm Optimisation (PSO) to minimise energy functions [1]. These methods, while interpretable, were computationally prohibitive for large proteins.

2.1.1 The Neural Network Revolution

The integration of Neural Networks (NNs) transformed TFM. Early deep learning approaches, such as DeepCNF and PSIPRED, focused on predicting 1D structural features (secondary structure) or 2D maps (contact/distance maps) to constrain the conformational search space [1]. The breakthrough came with the realisation that these constraints could be learned directly from the co-evolutionary signals embedded in MSAs.

AlphaFold2 (AF2) represented a paradigm shift by introducing an end-to-end differentiable architecture. Its Evoformer module utilises row and column attention to process MSA and pair representations iteratively, allowing the network to “reason” about spatial relationships before generating coordinates [1, 2]. This architecture essentially solved the static folding problem for single domains.

AlphaFold 3 (AF3) has recently advanced this further by replacing the Evoformer with a simpler “Pairformer” and adopting a diffusion-based generative module [2]. Unlike AF2, which builds the structure residue-by-residue, AF3 starts with a cloud of noise and progressively “denoises” it into a protein structure. This allows AF3 to model a broader range of biomolecules, including nucleic acids, ligands, and ions, with unprecedented accuracy [2, 3].

Insight: While AF3 improves interaction modelling, the shift to diffusion comes at a cost. The model no longer optimises a physical energy function but rather learns a statistical distribution of structures. This leads to issues such as stereochemical violations (chirality errors, clashing atoms) and “hallucinations” in disordered regions, where the model forces order upon intrinsically flexible segments [2].

2.1.2 The Single-Sequence Frontier

A critical limitation of AF2/AF3 is their dependence on deep MSAs. For “orphan proteins”—those with no evolutionary relatives—accuracy drops precipitously. RaptorX-Single and similar models (OmegaFold, ESMFold) address this by utilising Protein Language Models (PLMs) [8]. PLMs like ESM-1b and ProtTrans are trained on billions of unlabelled sequences (UniRef), learning the “grammar” of protein sequences. These embeddings capture implicit structural information, allowing RaptorX-Single to predict structures for orphan proteins and antibody CDR3 loops (which are highly variable and lack MSAs) with greater accuracy than MSA-based methods [8].

Implication: The success of single-sequence models suggests that the rules of folding are learnable from sequence alone, provided the model has seen enough diversity. This opens the door to designing synthetic proteins that have no natural homologs.

2.2 The Challenge of Thermodynamic Stability

Understanding how mutations affect that structure’s stability is equally vital as predicting the native structure. The prediction of Gibbs free energy change ($\Delta\Delta G$) has afflicted the field due to insufficiency of data and bias.

2.2.1 The Bias of Destabilisation

Experimental datasets for $\Delta\Delta G$ are heavily skewed toward destabilising mutations (since most random mutations break things). Therefore, naive ML models learn to predict “destabilisation” for almost any input. DDMut introduced a rigorous solution using Siamese Networks [9]. By inputting both the forward mutation ($A \rightarrow B$) and the reverse mutation ($B \rightarrow A$) and enforcing an antisymmetry loss ($\Delta\Delta G_{AB} = -\Delta\Delta G_{BA}$), DDMut forces the model to learn the physical reversibility of the thermodynamic state function. This approach, combined with graph-based representations of the local environment, yields state-of-the-art performance [9].

2.2.2 The Epistasis Debate

A profound theoretical question is whether residues evolve independently or if “epistasis” (interactions between residues) is the primary driver of stability. Recent studies exploring evolutionary models have yielded a surprising result: simple independent-site models often perform as well as complex epistatic models (like Potts models) in predicting mutational stability [5].

Insight: This suggests that for stability (unlike structure), the local conservation of a residue is a stronger signal than its coupling to distant residues. This has massive implications for model design: we may not need computationally expensive pairwise attention modules (like the Evoformer) to predict stability; lighter, site-specific models might suffice if trained on deep enough alignments.

2.2.3 Ultrafast Screening

The need for proteome-wide stability scanning has driven the development of ultrafast models like Pythia [18]. Unlike supervised models that require scarce $\Delta\Delta G$ labels, Pythia uses self-supervised learning on protein structures to learn a “micro-environment” representation. It can predict 50,000 mutations per second, enabling the mapping of entire evolutionary landscapes [18]. However, Pythia’s reliance on 3D structure means it fails if the initial structure is poor (low pLDDT), creating a dependency loop with structure predictors.

Other approaches include CART-based methods that classify mutations based on secondary structure and solvent accessibility [10], and unified frameworks like UniMutStab that handle insertions, deletions, and multiple mutations using shared-weight graph networks [11].

2.3 The Unstructured Proteome: Disorder and Dynamics

Approximately 30-40% of the eukaryotic proteome is intrinsically disordered. These Intrinsically Disordered Regions (IDRs) exist as dynamic ensembles but do not fold into stable 3D structures.

2.3.1 The “Unknown Significance” of IDRs

Mutations in IDRs are often categorized as “Variants of Unknown Significance” (VUS) because they do not disrupt a defined fold. However, recent research shows that single-point mutations

in IDRs can drastically alter the ensemble dimensions (radius of gyration) [16]. An analysis of ClinVar variants revealed that 42% of pathogenic mutations in IDRs cause the ensemble to compact, while 36% cause it to expand [16].

Implication: Pathogenicity in IDRs is driven by global conformational shifts—changes in the “fluffiness” or “compactness” of the chain—rather than local steric clashes. Current structure predictors like AF3, which attempt to predict a single static coordinate set, fundamentally misrepresent this biology. They often predict a compact “hallucination” for an IDR, obscuring the true dynamic defect caused by a mutation [2].

2.3.2 Interaction Networks

Mutations often manifest their effects by disrupting Protein-Protein Interactions (PPIs) [13] is demonstrated by the Contextual Graph Model (CGM). By looking at the “social network” of the protein, CGM predicts disease associations by integrating BERT-based language models with interaction graphs. The idea that interaction, stability, and structure are coupled is reinforced by this; a mutation might not unfold a protein but might prevent it from binding its partner.

For representation learning, identifying critical structural motifs that drive function, protein clustering is leveraged in advanced approaches. [17].

2.4 Sequence-to-Function Generalisation

“fragile generalisation” is a critical failure mode of deep learning in biology. When tested on synthetic or distantly related sequences; models trained on naturally occurring sequences frequently underperform.

2.4.1 Mechanistic Features vs. One-Hot Encoding

Standard CNNs/Transformers trained on raw sequence (one-hot encoding) collapse on novel test sets has been showed by a pivotal study on protein expression prediction [6]. Although, explicitly adding mechanistic features—biological priors like mRNA folding energy, codon adaptation index (CAI), and local charge profiles—via a “Geometric Stacking” approach significantly improved robustness [6].

Insight: Deep learning models are lazy; they latch onto local sequence motifs (e.g., “GC content at position 10”) that correlate with the training data but are not causal. Mechanistic features force the model to respect biological laws (e.g., “stronger mRNA folding inhibits translation initiation”), enabling it to generalise to sequences that violate the local motifs of the training set.

2.5 The Proteomics Data Ecosystem

Ultimately, the advancement of ML is bound to data quality. The field has suffered from a lack of standardisation in mass spectrometry (MS) proteomics.

2.5.1 The Scaling Laws of Proteomics

A recent multi-species benchmark containing 2.8 million Peptide-Spectrum Matches (PSMs) displayed a clear scaling law: dataset size is the primary driver of accuracy for peptide retention time prediction [12]. Models trained on this massive, curated dataset outperformed complex architectures trained on smaller, fragmented datasets [12].

Warning: The study also highlighted the dangers of Multi-Task Learning (MTL). Contrary to popular belief, blindly combining tasks (e.g., Retention Time + Collisional Cross Section) can degrade performance if the datasets are imbalanced or the tasks are physically discordant [7].

This suggests that “transfer learning” (pre-training on a large task, fine-tuning on a small one) is a safer strategy than simultaneous MTL.

2.6 Summary of Gaps and Opportunities

The literature reveals a convergent opportunity. We have powerful static structure predictors (AF3) that fail on dynamics [2, 3]. We have rigorous stability predictors (DDMut) that require pre-existing structures [9]. We have PLMs that handle orphan proteins but lack physical grounding [8]. And we have a clear signal that “mechanistic” features improve generalisation [6].

The missing link is a Unified Framework that:

- Predicts **Ensembles** (not just structures) to handle IDRs.
- Integrates **Mechanistic Features** to ensure generalisation.
- Uses **Antisymmetric Learning** for unbiased stability prediction.
- Operates on **Single Sequences** (via PLMs) to solve the orphan problem.

Our methodology, described below, aims to construct this unified framework.

3 Methodology

In response to the limitations identified in the literature, we propose the **Holo-GNN** (Holographic Geometric Neural Network). This framework is designed to be a comprehensive system for predicting protein properties, moving beyond the static coordinate generation of AlphaFold to a dynamic, ensemble-aware, and thermodynamically consistent model.

The name “Holographic” refers to the model’s ability to reconstruct the “full 3D image” and dynamic behaviour of the protein from the 1D information encoded in the sequence and mechanistic features, much like a hologram reconstructs a 3D light field from a 2D interference pattern.

3.1 Design Philosophy: Geometric Stacking

Our core methodological innovation is “Geometric Stacking” [6]. Standard Graph Neural Networks (GCNs) treat proteins as graphs where nodes are amino acids and edges are defined by distance ($< 8\text{\AA}$). We argue this is insufficient because it captures only geometric proximity, ignoring functional or evolutionary coupling.

In Holo-GNN, the graph topology is constructed using two distinct sets of edges:

- **Geometric Edges:** Defined by predicted spatial proximity (derived from an initial PLM contact map).
- **Mechanistic Edges:** Defined by biological similarity. For example, residues that share extreme codon usage bias or similar mRNA folding constraints are connected, even if they are distant in 3D space. This allows the network to learn “long-range” dependencies that are driven by translational kinetics rather than just folding energetics [6].

3.2 Data Acquisition and Provenance

A critical requirement for this study is the curation of a high-quality, standardised dataset. Based on the “multi-species benchmark” and “data utility” findings [7, 12], we reject the use of small, ad-hoc datasets in favour of massive, community-validated repositories.

3.2.1 Dataset Sources

The training corpus is assembled from the following specific sources. Note that all data retrieval is performed programmatically via public APIs to ensure reproducibility (see Table 2).

3.2.2 Data Pre-processing Protocol

To avoid the “fragile generalisation” and “data leakage” issues highlighted in the review [6, 12], we implement a rigorous **Peptide-Disjoint Splitting** strategy:

- **Sequence Clustering:** All sequences are clustered at 90% identity using MMseqs2.
- **Split Enforcement:** Entire clusters are assigned to Train, Validation, or Test. This ensures that no peptide in the test set has a near-neighbour in the training set.
- **Gold Standard Filtering:** For the Proteomics data, we only use PSMs with a False Discovery Rate (FDR) < 0.01 . For Stability data, we remove inconsistent entries where multiple experiments report conflicting $\Delta\Delta G$ values for the same mutation.

3.3 The Holo-GNN Architecture

The model consists of three interconnected modules: The Sequence Embedding Module, the Dual-Track Graph Module, and the Functional Output Heads.

3.3.1 Module 1: Sequence Embedding (The PLM Layer)

We utilise ESM-2 (3B parameter version) as the foundation [8].

- **Input:** Raw amino acid sequence $S = \{a_1, a_2, \dots, a_L\}$.
- **Processing:** The sequence is passed through the ESM-2 transformer layers.
- **Output:** A dense embedding vector $h_i \in \mathbb{R}^{1280}$ for each residue. This embedding captures the evolutionary context (homology) implicitly, solving the “Orphan Protein” issue without requiring explicit MSAs [1, 8].

3.3.2 Module 2: Mechanistic Feature Injection

Parallel to the PLM, we compute the Mechanistic Features [6]:

- mRNA_{fold} : The minimal free energy (MFE) of the mRNA window corresponding to residues $i - 5$ to $i + 5$.
- CAI_i : The Codon Adaptation Index for the codon at position i .
- Charge_i : Local charge density.

These features are concatenated with the PLM embedding: $h'_i = [h_i || \text{mRNA}_{fold} || \text{CAI}_i || \text{Charge}_i]$.

3.3.3 Module 3: Dual-Track Graph Neural Network

We construct the graph $\mathcal{G} = (\mathcal{V}, \mathcal{E})$ where nodes $\mathcal{V}$ are residues with features h'_i . We draw inspiration from graph-based clustering approaches that identify critical structural motifs [17].

Edge Construction:

- E_{geo} : Connect i, j if contact probability $P(i, j) > 0.5$ (predicted by a lightweight pairwise head).

- E_{mech} : Connect i, j if $|CAI_i - CAI_j| < \epsilon$ AND $|mRNA_i - mRNA_j| < \delta$. This links residues that are “translationally coupled.”

Message Passing: We employ a Graph Attention Network (GAT) layer that allows nodes to aggregate information from both geometric and mechanistic neighbours. This fusing of spatial and biological information is the core implementation of “Geometric Stacking” [6].

3.3.4 Module 4: The Siamese Stability Head

To predict $\Delta\Delta G$, we use the Siamese Architecture derived from DDMut [9].

- **Input:** Two graphs, $\mathcal{G}_{WT}$ (Wild Type) and $\mathcal{G}_{MT}$ (Mutant).
- **Shared Encoder:** Both graphs pass through the same Holo-GNN encoder to produce graph-level embeddings Z_{WT} and Z_{MT} .
- **Prediction:** A Multilayer Perceptron (MLP) predicts $\Delta\Delta G = \text{MLP}(Z_{MT} - Z_{WT})$.
- **Antisymmetry Loss:** We explicitly minimise:

$$\mathcal{L}_{antisym} = (\Delta\Delta G_{WT \rightarrow MT} + \Delta\Delta G_{MT \rightarrow WT})^2 + (\Delta\Delta G_{pred} - \Delta\Delta G_{exp})^2$$

This guarantees that the model respects the thermodynamic law that the free energy change of a reverse mutation is the negative of the forward mutation. We also incorporate insights from UniMutStab for handling insertions and deletions [11].

3.3.5 Module 5: The Ensemble Dimension Head (IDR)

To handle IDRs, we avoid predicting a single coordinate set. Instead, for residues classified as “disordered” (using a dedicated classifier token), we predict the Radius of Gyration (R_g) and End-to-End Distance (R_{ee}) as probability distributions (Gaussian).

- **Output:** μ_{Rg}, σ_{Rg} .
- **Loss:** Kullback-Leibler (KL) divergence against the experimental distributions derived from the ClinVar analysis [16]. This allows the model to predict if a mutation causes “Compaction” or “Expansion” of the ensemble, a key driver of IDR pathogenicity.

3.4 Training Strategy

We employ a **Curriculum Learning** approach to handle the diverse data tasks.

1. **Phase 1: Self-Supervised Geometry (Warm-up).** The model is trained on UniRef sequences to predict masked residues (MLM) and pairwise contacts (self-supervised structure learning). This primes the PLM and Graph modules [8].
2. **Phase 2: Transfer Learning for Stability.** We freeze the PLM layers and train the GNN + Siamese Head on stability datasets. This ensures the limited $\Delta\Delta G$ data is used efficiently to fine-tune the high-level representations [9, 18].
3. **Phase 3: Multi-Task Fine-Tuning.** We unfreeze all layers and train simultaneously on Proteomics (RT/CCS) and Pathogenicity tasks. We use Gradient Normalisation (Grad-Norm) to balance the learning rates of different tasks, preventing the larger proteomics dataset from drowning out the stability signals. Based on insights from proteomics ML research, we prioritise transfer learning over naive multi-task learning [7].

3.5 Contributions

Our research provides the following distinct contributions to the field:

1. **A Unified Solution for Static and Dynamic Proteomics:** By integrating the “Ensemble Dimension” head, we provide the first deep learning framework capable of explicitly modelling the conformational expansion/compaction of Intrinsically Disordered Regions, addressing a critical blind spot of AlphaFold [2, 16].
2. **Resolution of the “Fragile Generalisation” Problem:** We demonstrate that the injection of “Mechanistic Features” (mRNA, Codon Bias) via Geometric Stacking allows the model to generalise to sequences with novel evolutionary histories, outperforming raw-sequence models in synthetic biology tasks [6].
3. **An Unbiased, MSA-Free Stability Predictor:** By combining PLMs (which handle orphans [8]) with Siamese Networks (which remove bias [9]), we create a stability predictor that is both physically rigorous and applicable to the “dark matter” of the protein universe (orphan proteins).
4. **A Standardised Benchmarking Protocol:** We release the re-processed, peptide-disjoint train/test splits for the 2.8M PSM proteomics dataset, providing the community with a definitive “Gold Standard” for future ML model comparison [12].

Table 1: Detailed Architectural Comparison

Feature	AlphaFold 3 (AF3)	DDMut	RaptorX-Single	Holo-GNN (Ours)
Primary Input	MSA (Evolutionary) + Templates	3D Structure (PDB)	Single Sequence (PLM)	Single Sequence + Mechanistic Features
Core Architecture	Pairformer + Diffusion	Graph Neural Network (GNN)	Transformer + Evoformer	Dual-Track GAT + Siamese Network
Stability Logic	Indirect (pLDDT confidence)	Antisymmetric (Siamese)	None / Indirect	Antisymmetric (Siamese)
IDR Handling	Static “Hallucination”	N/A	Static	Ensemble Distribution (μ, σ)
Orphan Capability	Poor (Requires MSA)	N/A (Requires Structure)	Excellent	Excellent (PLM-driven)
Indel Support	Limited	Single Point Focus	Limited	Full Support (Elastic Graph)

Table 2: Dataset Inventory and Specifications

Dataset Name	Size / Scale	Features Extracted	Specific Use Case
MassIVE-KB	2.8 Million PSMs	Sequence, Retention Time, CCS	Training the Proteomics Head; Validating Generalisation.
FireProtDB	~5,000 Verified Mutations	$\Delta\Delta G$ (kcal/mol), pH, Temp	Fine-tuning the Siamese Stability Head.
Mega-scale cDNA	780,000 Mutations	Proteolysis rates	Large-scale pre-training of stability features (UniMutStab approach).
ClinVar (IDR)	42,000 IDR Variants	Pathogenicity (Benign/Pathogenic)	Calibrating the Ensemble Dimension Head (Compaction vs Expansion).
UniRef50	50 Million Clusters	Sequence only	Self-supervised pre-training (Masked Language Modelling).

4 Benchmark analysis

4.1 Background

Protein structure prediction has undergone a revolutionary transformation with the advent of deep learning-based methods. The accurate prediction of three-dimensional protein structures from amino acid sequences is fundamental to understanding biological function, drug discovery, and protein engineering.

4.2 Methods Under Evaluation

This benchmark evaluates five cutting-edge prediction methods:

1. **AlphaFold2** (DeepMind): The breakthrough method that achieved near-experimental accuracy in CASP14
2. **ESMFold** (Meta AI): A language model-based approach that eliminates MSA dependency
3. **BioEmu** (Conformational Sampler): A method for exploring protein conformational landscapes
4. **Boltz1** (Diffusion-based): An advanced diffusion model for structure generation
5. **RoseTTAFold2** (Baker Lab): A three-track neural network architecture

Table 3: Test Sequences Used in Benchmark

ID	Method(s)	Length	Sequence
test_a5e17	AF2, RF2, Boltz1	59	PIAQIHILEGRSDEQKETLIREVSEAIS RSLDAPLTSVRVIITEMAKGHFGIGGELASK
test_518d3	ESMFold	68	(Sequence from ESMFold test)
calmodulin_b5361	BioEmu	148	MADQLTEEQIAEFKEAFSLFDKDGDTITT KELGTVMRSLGQNPTEAELQDMINEVDAD GNGTIDFPEFLTMMARKMKDTSSEEEIRE AFRVFDKDGNGYISAAELRHVMTNLGEKL TDEEVDEMIREADIDGDGQVNYEEFVQMMTAK

4.3 Test Sequences

5 Methods Overview

5.1 AlphaFold2

5.1.1 Architecture

AlphaFold2 employs a sophisticated architecture consisting of:

- **Evoformer:** Processes MSA and pairwise representations through 48 blocks
- **Structure Module:** Converts representations to 3D coordinates
- **Recycling:** Iteratively refines predictions through multiple passes

5.1.2 Configuration Used

Table 4: AlphaFold2 Configuration Parameters

Parameter	Value
Version	1.5.5 (ColabFold)
Model Type	alphafold2_ptm
Number of Models	5
Number of Recycles	3
MSA Mode	mmseqs2_uniref_env
Max Sequence Depth	512
Max Extra Sequences	5120
Use Dropout	True
Use Templates	False
Number of Relaxation Steps	0
Pair Mode	unpaired_paired
Pairing Strategy	greedy
Rank By	pLDDT
Random Seed	0
Use bfloat16	True

5.2 ESMFold

5.2.1 Architecture

ESMFold leverages the ESM-2 protein language model:

- **ESM-2 Backbone:** 15 billion parameter language model
- **Folding Module:** Converts embeddings to structure
- **No MSA Required:** Single-sequence prediction capability

5.2.2 Configuration Used

Table 5: ESMFold Configuration Parameters

Parameter	Value
Recycles	3 (r3)
Mode	default
pTM Score	0.817

5.3 BioEmu

5.3.1 Architecture

BioEmu is designed for conformational sampling:

- **Generative Model:** Samples multiple conformational states
- **Molecular Dynamics Integration:** Optional refinement with OpenMM
- **Clustering:** Foldseek-based structural clustering

5.4 Boltz1

5.4.1 Architecture

Boltz1 employs a diffusion-based approach:

- **Pairformer:** 64 transformer blocks for pairwise representations
- **Diffusion Module:** Iterative denoising for structure generation
- **Multi-entity Support:** Handles proteins, ligands, and complexes

5.4.2 Configuration Used

5.5 RoseTTAFold2

5.5.1 Architecture

RoseTTAFold2 uses a three-track architecture:

- **1D Track:** Sequence representations
- **2D Track:** Pairwise distance and orientation features
- **3D Track:** Coordinate predictions

Table 6: Boltz1 Hyperparameters

Parameter	Value
<i>Architecture</i>	
Token Dimension (s)	384
Pair Dimension (z)	128
Number of Bins	64
<i>Pairformer</i>	
Number of Blocks	64
Number of Heads	16
Dropout	0.0
<i>Diffusion</i>	
Sigma Minimum	0.0001
Sigma Maximum	160.0
Sigma Data	16.0
P Mean	-1.2
P Std	1.5
Validation Sampling Steps	200
<i>Training</i>	
Recycling Steps	3
Max Learning Rate	0.001
Weight Decay	0.003

Table 7: RoseTTAFold2 Configuration Parameters

Parameter	Value
Method	RoseTTAFold2
Parameters	RF2_apr23
Number of Recycles	6
MSA Method	mmseqs2
Number of Models	1
Sequence Length	59

5.5.2 Configuration Used

6 Results

6.1 AlphaFold2 Detailed Results

6.1.1 Model Rankings

AlphaFold2 generated predictions from 5 different PTM models, ranked by pLDDT score:


Table 8: AlphaFold2 Model Rankings and Quality Metrics

Rank	Model	pTM	Max PAE	Mean pLDDT	Min pLDDT
1	model_3	0.78	17.89	97.2	82.88
2	model_4	0.77	17.95	96.8	82.69
3	model_2	0.75	20.61	95.9	71.88
4	model_5	0.77	21.50	95.7	67.44
5	model_1	0.74	21.16	95.3	75.19

Table 9: Per-Residue pLDDT Scores for AlphaFold2 Best Model (model_3)

Position	1-10	11-20	21-30	31-40	41-50	51-59
Range	96.0-98.6	97.2-98.7	98.4-98.8	98.2-98.7	95.1-98.6	82.9-97.2
Quality	Very High	Very High	Excellent	Excellent	High	Moderate

6.1.2 Per-Residue pLDDT Analysis (Best Model - Rank 1)



6.1.3 Predicted Aligned Error (PAE) Analysis

The PAE matrix for AlphaFold2’s best model shows:

- **Diagonal Region:** Very low PAE (0.25 Å) indicating high self-confidence
- **Core Domain** (residues 10-45): Low inter-residue PAE (< 3 Å)
- **C-terminal Region** (residues 50-59): Higher uncertainty (up to 17.89 Å)
- **Global Maximum PAE:** 17.89 Å

6.2 ESMFold Results

6.2.1 Quality Metrics

6.2.2 PAE Matrix Characteristics

The 68×68 PAE matrix from ESMFold reveals:

- **N-terminal Region (1-20):** Extremely low PAE (0.001-2.0 Å) - highest confidence
- **Central Region (20-45):** Low to moderate PAE (2-8 Å)

Table 10: ESMFold Quality Metrics

Metric	Value
pTM Score	0.817
Sequence Length	68 residues
Recycles	3
Mode	default

- **C-terminal Region (45-68):** Moderate to high PAE (8-20 Å)

Table 11: ESMFold PAE Regional Summary

Region	Residues	PAE Range (Å)	Confidence
N-terminus	1-20	0.001 - 2.0	Very High
Core	20-45	2.0 - 8.0	High
C-terminus	45-68	8.0 - 20.0	Moderate

6.3 BioEmu Conformational Sampling Results

6.3.1 Target: Calmodulin

Calmodulin is a calcium-binding messenger protein known for its conformational flexibility, making it an ideal target for conformational sampling.

Sequence (148 residues):


MADQLTEEQIAEFKEAFSLFDKDGDTITTKELGTVMRSLGQNPTEAELQDMINEVDAD
 GNGTIDFPEFLTMMARKMKDSTDSEEEIREAFRVFDKDGNGYISAAELRHVMTNLGEKLT
 DEEVDEMIREADIDGDGQVNYEEFVQMMTAK

6.3.2 Clustering Results

BioEmu generated 10 conformational samples, which were clustered using Foldseek into 4 distinct structural groups:

Table 12: BioEmu Foldseek Clustering Results

Cluster	Count	Sample IDs	Interpretation
0	3	0, 2, 8	Alternative conformational state
1	5	1, 3, 4, 6, 9	Dominant state (50%)
2	1	5	Rare/unique conformation
3	1	7	Rare/unique conformation



6.4 Boltz1 Results

6.4.1 Model Architecture Analysis

Boltz1 represents one of the most sophisticated architectures tested:

Table 13: Boltz1 Architecture Comparison with AlphaFold2

Component	Boltz1	AlphaFold2
Transformer Blocks	64	48
Token Dimension	384	256
Pair Dimension	128	128
Attention Heads	16	8
Structure Generation	Diffusion	Direct
Ligand Support	Yes	Limited

6.4.2 Diffusion Parameters

The diffusion process in Boltz1 is characterized by:

$$\sigma(t) = \sigma_{\min}^{1-t} \cdot \sigma_{\max}^t, \quad t \in [0, 1] \quad (1)$$

Where:

- $\sigma_{\min} = 0.0001$ (minimum noise level)
- $\sigma_{\max} = 160.0$ (maximum noise level)
- $\sigma_{\text{data}} = 16.0$ (data distribution spread)

6.5 RoseTTAFold2 Results

6.5.1 Configuration Details

6.5.2 MSA Generation

RoseTTAFold2 utilized comprehensive MSA databases:

- `uniref.a3m`: UniRef sequence database
- `bfd.mgnify30.metaeuk30.smag30.a3m`: Combined environmental databases

Table 14: RoseTTAFold2 Run Parameters

Parameter	Value
Method	RoseTTAFold2
Parameter Set	RF2_apr23
Recycles	6 (highest among tested methods)
MSA Method	mmseqs2
Number of Models	1
Test ID	test_X1_a5e17

- pdb70.m8: Template database hits

7 Comparative Analysis

7.1 pTM Score Comparison

Table 15: pTM Score Comparison Across Methods

Method	pTM Score	Rank
ESMFold	0.817	1
AlphaFold2 (best)	0.780	2
AlphaFold2 (model_4)	0.770	3
AlphaFold2 (model_5)	0.770	3
AlphaFold2 (model_2)	0.750	5
AlphaFold2 (model_1)	0.740	6




Table 16: Comprehensive Method Comparison

Feature	AF2	ESMFold	BioEmu	Boltz1	RF2
MSA Required	Yes	No	Variable	Yes	Yes
Number of Models	5	1	N/A	1	1
Recycles	3	3	N/A	3	6
Conformational Sampling	No	No	Yes	No	No
Ligand Support	Limited	No	No	Yes	Limited
Speed	Moderate	Fast	Slow	Moderate	Moderate

Table 17: MSA Dependency and Data Sources

Method	MSA Required	Databases Used
AlphaFold2	Yes	UniRef, Environmental (BFD, MGnify, MetaEuk, SMAG)
ESMFold	No	None (PLM embeddings only)
BioEmu	Variable	Depends on configuration
Boltz1	Yes	UniRef via MMseqs2
RoseTTAFold2	Yes	UniRef, BFD, MGnify, PDB70

Table 18: Computational Resource Comparison

Method	GPU Memory	Inference Time	Scalability
AlphaFold2	High	Moderate	Good
ESMFold	Moderate	Fast	Excellent
BioEmu	High	Slow (sampling)	Moderate
Boltz1	Very High	Moderate	Good
RoseTTAFold2	High	Moderate	Good

7.2 Method Characteristics Comparison

7.3 MSA Dependency Analysis

7.4 Computational Considerations

8 Discussion

8.1 Key Findings

8.1.1 Best Overall Performance

ESMFold achieved the highest pTM score (0.817) among all tested methods, demonstrating that language model-based approaches can match or exceed traditional MSA-dependent methods for single-chain predictions.

8.1.2 Most Robust Ensemble

AlphaFold2 provides the most comprehensive prediction pipeline with:

- 5 different model variants for ensemble averaging
- Detailed per-residue confidence scores (pLDDT)
- Comprehensive PAE matrices for domain analysis

- Well-calibrated uncertainty estimates

8.1.3 Conformational Insights

BioEmu uniquely captures conformational diversity:

- Generated 4 distinct structural clusters from 10 samples
- Identified dominant (50%) and rare conformational states
- Particularly valuable for flexible proteins like Calmodulin

8.1.4 Advanced Architecture

Boltz1 represents the frontier of structure prediction:

- 64 pairformer blocks (vs. 48 in AF2)
- Diffusion-based generation allows for uncertainty sampling
- Native support for ligands and multi-chain complexes

8.1.5 Deep Refinement

RoseTTAFold2 with 6 recycles (highest tested) provides:

- More iterations for structure refinement
- Three-track architecture for comprehensive feature integration
- Strong performance on protein complexes

8.2 Confidence Score Interpretation

8.2.1 pLDDT Score Guidelines

Table 19: pLDDT Score Interpretation

pLDDT Range	Confidence	Interpretation
> 90	Very High	Backbone likely accurate within 1-2 Å
70-90	High	Backbone generally reliable
50-70	Low	Significant uncertainty, use caution
< 50	Very Low	Likely disordered or incorrectly predicted

8.2.2 PAE Score Guidelines

Table 20: PAE Score Interpretation

PAE Range (Å)	Confidence	Interpretation
0-5	Very High	Residue positions well-defined relative to each other
5-10	High	Good relative positioning
10-20	Moderate	Some uncertainty in relative positions
> 20	Low	Poor confidence in relative positions

8.3 Method Selection Guidelines

1. **For rapid single-chain predictions:** Use **ESMFold**

- No MSA computation required
- Highest pTM score in this benchmark
- Fastest inference time

2. **For robust ensemble predictions:** Use **AlphaFold2**

- Multiple model variants for averaging
- Well-calibrated confidence metrics
- Extensive validation in CASP

3. **For conformational dynamics:** Use **BioEmu**

- Samples multiple structural states
- Captures rare conformations
- Ideal for flexible proteins

4. **For protein-ligand complexes:** Use **Boltz1**

- Native ligand support
- Advanced diffusion architecture
- Multi-chain capability

5. **For protein complexes:** Use **RoseTTAFold2**

- Deep recycling (6 iterations)
- Three-track architecture
- Strong multimer performance

9 Conclusion

This comprehensive benchmark analysis of five state-of-the-art protein structure prediction methods reveals distinct strengths and optimal use cases for each approach:

1. **ESMFold** achieves the highest pTM score (0.817) while eliminating MSA dependency, making it ideal for rapid predictions.
2. **AlphaFold2** provides the most robust prediction pipeline with well-calibrated confidence metrics across 5 model variants.
3. **BioEmu** uniquely captures protein dynamics through conformational sampling, generating 4 distinct structural clusters for Calmodulin.
4. **Boltz1** represents the architectural frontier with 64 pairformer blocks and native ligand support through diffusion-based generation.
5. **RoseTTAFold2** offers deep refinement with 6 recycles and strong performance on protein complexes.

The choice of method should be guided by the specific research question: speed vs. accuracy, single chain vs. complex, static structure vs. conformational ensemble.

Table 21: Benchmark Summary - Best Performers by Metric

Metric	Best Performer	Value
pTM Score	ESMFold	0.817
Mean pLDDT	AlphaFold2 (model_3)	97.2
Lowest Max PAE	AlphaFold2 (model_3)	17.89 Å
Model Ensemble Size	AlphaFold2	5 models
Conformational Coverage	BioEmu	4 clusters
Architectural Depth	Boltz1	64 blocks
Recycle Depth	RoseTTAFold2	6 recycles
MSA-Free Prediction	ESMFold	Yes
Ligand Support	Boltz1	Native

10 Summary Statistics

A Glossary

pTM Predicted Template Modeling score - global structure quality metric (0-1)

pLDDT Predicted Local Distance Difference Test - per-residue confidence (0-100)

PAE Predicted Aligned Error - pairwise residue position uncertainty (Å)

MSA Multiple Sequence Alignment - evolutionary sequence information

Evoformer AlphaFold2’s main transformer block processing MSA and pair representations

Pairformer Boltz1’s transformer architecture for pairwise features

Recycles Iterative refinement passes through the network

Diffusion Generative modeling approach using iterative denoising

B References

1. Jumper, J., et al. (2021). Highly accurate protein structure prediction with AlphaFold. *Nature*, 596(7873), 583-589.
2. Lin, Z., et al. (2023). Evolutionary-scale prediction of atomic-level protein structure with a language model. *Science*, 379(6637), 1123-1130.
3. Baek, M., et al. (2021). Accurate prediction of protein structures and interactions using a three-track neural network. *Science*, 373(6557), 871-876.
4. Mirdita, M., et al. (2022). ColabFold: making protein folding accessible to all. *Nature Methods*, 19(6), 679-682.
5. Van Kempen, M., et al. (2023). Fast and accurate protein structure search with Foldseek. *Nature Biotechnology*, 42, 243-246.

C Conclusion

The transition from the “Structure Prediction Era” to the “Dynamic Function Era” requires a fundamental rethinking of how we model proteins. We can no longer rely solely on the static snapshots provided by crystallography or the hallucinations of diffusion models [2, 3]. By acknowledging the thermodynamic reality of proteins—that they are ensembles defined by energy landscapes, not just coordinates [16]—and by respecting the “mechanistic” rules of their biogenesis (transcription and translation) [6], we can build models that are not only more accurate but more biologically grounded.

Holo-GNN represents a first step toward this unified vision. By stacking geometric deep learning on top of evolutionary language models [17] and enforcing physical symmetries [9], we offer a robust tool for the next generation of protein engineering and precision medicine. Our approach addresses the key limitations identified across the literature: static prediction [3], poor generalisation [6], MSA dependency [8], stability prediction bias [9], and disordered region misrepresentation [16].

References

- [1] Artificial intelligence-based methods for protein structure prediction: a survey.
- [2] Accurate structure prediction of biomolecular interactions with AlphaFold 3.
- [3] AlphaFold 3: an unprecedented opportunity for fundamental research and drug development.
- [4] Using AlphaFold to predict the impact of single mutations on protein stability and function.
- [5] Exploring Evolution to Uncover Insights Into Protein Mutational Stability.
- [6] Shen et al. (2025). Improving the Generalisation of Protein Expression Models. *Nucleic Acids Research*.
- [7] Dens et al. (2024). ML Strategies to Tackle Data Challenges in Proteomics. *bioRxiv*.
- [8] Jing et al. (2024). Single-Sequence Protein Structure Prediction (RaptorX-Single). *PNAS*.
- [9] DDMut: Predicting effects of mutations on protein stability using deep learning.
- [10] Prediction of protein mutant stability using Classification and Regression Tool (CART).
- [11] Shared-weight graph framework for comprehensive protein stability prediction across diverse mutation types (UniMutStab).
- [12] A multi-species benchmark for training and validating mass spectrometry proteomics machine learning models.
- [13] Predicting mutation-disease associations through protein interactions via deep learning.
- [14] Genetic variant classification by predicted protein structure: A case study on IRF6.
- [15] Integrating Artificial Intelligence and Bioinformatics Methods to Identify Disruptive STAT1 Variants Impacting Protein Stability and Function.
- [16] Single-point mutations in disordered proteins.
- [17] Clustering for Protein Representation Learning.
- [18] Structure-based self-supervised learning enables ultrafast protein stability prediction upon mutation (Pythia).